

2026 NAIROBI ONSITE · TALK + WORKSHOP · 90 MINUTES


Measuring **agency** in a digital context

How can we tell when digital tool use is expanding or undermining human agency? What is agency? Why should we measure it? How can we measure it to know whether digital tools are actually supporting it?

Zezen Wu · Behavioral Scientist, The Agency Fund

WHERE THE AGENCY FUND STARTS

Our worldview of agency in a flywheel



AI can **expand** agency, and it can hamper it

CAN EXPAND

JIGSAW AT GOOGLE · ETHNOGRAPHY

Feelings of agency, not task completion, drive AI use: people stay with tools that make them feel more capable, even when the tools make mistakes. AI can strengthen agency in five ways: instrumental, cognitive, affective, relational, and structural.

60

advanced AI users interviewed

3

countries: US, Singapore, Germany

CAN HAMPER

BASTANI ET AL. · PNAS FIELD RCT

Answers on demand became a crutch: a ChatGPT-style math tutor lifted the practice scores of ~1,000 students, but once it was taken away they scored below peers who never had it. Hints-only guardrails erased the harm, and nobody felt the loss.

+48%

practice scores, AI in hand

-17%

exam scores, AI removed

CAN HAMPER

SHAW AND NAVE · 3 EXPERIMENTS

Cognitive surrender: offered an optional AI on reasoning questions, 1,372 participants consulted it on more than half of trials and adopted its answer, right or wrong. Accuracy tracked the tool, and confidence rose even when the answer was wrong.

+25 pts

accuracy when the AI was right

-15 pts

when it was wrong, vs brain-only

In all three studies, people looked engaged and felt capable while it happened. Usage and satisfaction cannot tell expansion from erosion, **which is why we measure agency directly.**

Sources: [How AI Is Affecting Human Agency](#) (Jigsaw, July 2025) · [Generative AI without guardrails can harm learning](#) (PNAS 122(26), 2025) · Thinking, Fast, Slow, and Artificial: the rise of cognitive surrender (Wharton working paper, 2026).

Agency needs to be defined to the context and measured as an outcome

For example, in **social cognitive theory**, people are agents who intentionally shape their circumstances. Agency has four core elements in this framework:

1 Intentionality

Forming plans and committing to act on them.

SUB-CONSTRUCT · GOAL COMMITMENT

The learner states a business goal in her own words and **commits to a concrete next step** before the session ends.

2 Forethought

Setting goals and anticipating the likely outcomes of actions.

SUB-CONSTRUCT · PLANNING AND TRADEOFFS

She **weighs two ideas against each other**, asking what each would cost, earn, and risk in her local market.

3 Self-reactiveness

Monitoring progress and adjusting course along the way.

SUB-CONSTRUCT · COURSE CORRECTION

He returns with progress updates and **revises the plan** when the first supplier falls through, instead of abandoning the goal.

4 Self-reflectiveness

Evaluating one's own thinking, capability, and results.

SUB-CONSTRUCT · CALIBRATED SELF-BELIEF

After the pitch she reviews what worked and what did not, and her **stated confidence tracks her real progress**.

[Bandura \(2001\), Social cognitive theory: an agentic perspective](#). This is one lens among several: Sen's capability approach and Jigsaw's five agency types (slide 3) name different sub-constructs. Whatever the framework, pick one whose constructs your product can actually observe.

Four types of evidence around an AI business mentor

Agency with respect to the end goals

Can users make progress toward outcomes they value, whether through, around, or beyond the tool?

Agency with respect to the digital tool

Can users understand, question, adapt, and control the tool, deciding when and how to rely on it?

Internally experienced agency

People's thoughts and feelings



ON-PLATFORM · IN-APP MICRO-SURVEY

Short in-app surveys probing self-efficacy or perceived control over the ultimate goal: how confident is the user about developing a successful business idea?



OFF-PLATFORM · PHONE OR IN-PERSON SURVEY

Longer surveys using validated agency measures, probing the same constructs systematically over time.



ON-PLATFORM · IN-APP MICRO-SURVEY

Short in-app surveys on whether users felt able to understand, question, or adopt the AI mentor's advice.



OFF-PLATFORM · VALIDATED SCALES

Validated measures of trust in AI and of perceived dependence versus empowerment when using the tool.

Externally expressed agency

People's choices and actions



ON-PLATFORM · CONVERSATION LOGS

Longitudinal signals of agency in conversation logs, especially setting and navigating longer-term goals, or resilience despite difficulties.



OFF-PLATFORM · RUBRIC-RATED PITCHES

Surveys tracking how business ideas develop; rubric-rated observations of ideas and pitches, ideally linked to outcomes like earnings.



ON-PLATFORM · DIGITAL TRACES

Interaction patterns and conversation logs showing whether users understand, question, or adopt the mentor's advice.




OFF-PLATFORM · OBSERVATION AND INTERVIEWS

Observation or interview data, such as usability studies, on how the tool gets used in practice.

In the workshop you will work in one corner of this table: externally expressed agency, read from conversation logs.

Agency measurement sits at **Level 3**



Nobody can run an RCT on every product change. **Validated agency measures at Level 3** move when the product changes, and they speak to what Level 4 cares about. That is what lets teams anchor their A/B tests on something that matters.

Teams should specify what agency means for their specific users and contexts

1 Define agency for your local problem

What kind of agency is the tool supposed to support? For whom? To what end? Use the two distinctions to move from a broad idea to a structured definition: agency with respect to the **end goals** versus the **digital tool**, and agency that is **internally experienced** versus **externally expressed**.

2 Measure with the right tools, on the right data, at the right moment

On-platform: in-app surveys, chatbot interactions, task behavior, and other digital traces. **Off-platform:** phone and in-person surveys, interviews, observations, administrative records. Self-reports are the most common tool and carry known biases; digital traces reveal what surveys miss, especially when linked to users' goals.

3 Test and validate your measurement tools iteratively

A measure is **reliable** if it produces consistent results under similar conditions. It has **construct validity** if it captures agency as you define it rather than a neighbor like mood or satisfaction, **content validity** if it covers the parts of agency that matter in your context, and **criterion validity** if its scores track outcomes you trust. For LLM-based measurement, compare the model's ratings against a gold-standard dataset coded by trained human reviewers using a clear rubric.

PART 2 · WORKSHOP

Now you find the agency, **line by line.**

Why conversations? Because they can hold rich signals of agency, especially when a human is talking with an AI. You will read 18 synthetic conversations across three domains, answer four questions and an open note, and compare your labels with an AI judge at the end.

ACTIVITY • ABOUT 40 MINUTES

Your turn: code conversations like a measurement team

1 Pick one card from a domain pile

Education, health, or agriculture: choose the world closest to your work. The card's QR code is your personal labeller link.

2 Open the link on your phone or laptop

You will see six conversations between users and a fictional AI product, plus the four coding questions.

3 Code alone, top to bottom, about 20 minutes

Answer the four questions for each conversation and leave one open note. Three or four conversations is a good pace; nobody needs to finish six.

4 Then compare labels at your table, about 15 minutes

Find one conversation where you disagreed. Argue it out: what evidence did each of you use? Nominate your hardest case for the debrief.

THREE ARCHETYPES, ALL FICTIONAL



AI educational mentor

Study and career support for secondary students



AI health assistant

Maternal and newborn health support



AI agricultural advisor

Advisory support for smallholder farmers

All synthetic. The products are fictional and every conversation was written for this exercise. No real user data, no real organization.

Four questions, one open note

1

INTENTIONALITY + FORETHOUGHT

Goals and plans

Does the user set or refine goals and plans of their own during the conversation, beyond what the AI suggests?

2

SELF-REACTIVENESS

Acting and adjusting

Does the user report taking action, tracking progress, or adjusting course when things change? First sessions give little evidence here; Mixed is a fair answer.

3

SELF-REFLECTIVENESS

Reflecting

Does the user reflect on their own experience and capabilities: what worked, what they can handle, what they believe now? A negative self-assessment still counts as reflection.

4

PERSONAL · PROXY · COLLECTIVE

Who is the author?

People act directly, through others (including an AI), and together with others. When this user leans on the AI or on people, do they stay the author of their choices, or hand the decision over?

OPEN NOTE

Anything else you notice, especially signs the tool may be undermining agency: over-reliance, decisions deferred to the AI, growing discouragement. Name the mode you see: personal, proxy, or collective.



SET

5 min

10 min

15 min

20 min

-1 min

+1 min

Start

Reset

Now we find out how much we disagreed

1 Human agreement, per domain

We open Calibrate's agreement view together: how often did coders give the same answer to the same question about the same conversation?

2 The LLM judge, same rubric

An AI evaluator has already run your exact four questions across all 18 conversations. We compare its verdicts with the room's.

3 The hard cases

Two or three tables tell us where they split and what evidence each side used. Those arguments are the raw material for a sharper rubric.

The point of the exercise: agency leaves signals in text, and we want to identify them **consistently**. Human coders set the standard. Once a rubric survives this contest with human judgment, an LLM can apply it across thousands of conversations. That pairing is what makes agency measurement scale.

Calibrate your expectations

When trained medical experts coded real conversations in our Noora Health pilot, they agreed with each other less than half the time. Disagreement is normal. It is also the material we work with.

This was Step 3 in miniature: humans first, the machine second, and the argument between them is where the measure improves.

Agency measurement is interesting, necessary, and complex. **Start anyway.**

The blog →

Measuring Agency in a Digital Context, by The Agency Fund and Google Jigsaw

The Claude skill →

ask-agency-measurement: a coach that turns these frameworks into a plan for your product

The NotebookLM →

Curated papers and validated agency scales you can interrogate conversationally

The evaluation playbook →

The full four-level framework for evaluating GenAI products, level by level